

# CS3210 Lab 2 Report


Wang Xiyu (A0282500R)

September 11, 2025

## Ex11 Performance Metrics

This sections shows the comparison of i7-7700(4 cores 8 threads) and Xeon Silver 4114 (10 cores 20 cores) in performance measured by instruction per cycle(IPC), wall clock time(elapsed time) and MFLOPS(Million floating point operations). Both processors have the same core architecture.


### Instruction per Cycle & Elapsed time



We can observe that for i7-7700 IPC moderately decreases from 1 to 4 threads, and IPC on Xeon 4114 also moderately decreases at similar rate from 1 to 10 threads, indicating that there exist overhead of thread scheduling, therefore the number of instructions for the actual task counted in floating point operations decreases. But the overall performance measured in wall clock time improves rapidly as there are more processing units being utilized.

When the number of threads = the number of PUs the wall clock time reaches minimum, as all PUs are utilized while and there is no extra overhead associated with scheduling incurs. When the number of PUs < number of threads < number of cores, instruction per cycle for both machines decreases more quickly, as other than the task allocation overhead, there is memory overhead incurred as well, since both PUs in the same core share the same cache, and when both PUs are being utilized, there can be memory contention.

### MFLOPS



For both processors we can observe that MFLOPS peaks when the number of threads = the number of cores, this is intuitive as there are more processing units are being utilized and the number of floating point operation increases. MFLOPS starts to decrease when number of cores  $\leq$  the number of threads  $\leq$  the number of processing units, this is due to overheads such as memory contention, as both PUs residing in the same core share the same LLC (last level cache), causing the overall

## Ex12&13 Optimization & Performance Analysis

As we can observe from the perf statistics, there are only scalar floating point operations being performed, which is a wasted opportunity for SIMD parallelism as we are accessing array elements and we can exploit the spatial locality if we can access contiguous cache in consecutive instructions. As 2D arrays are stored row-wise, we can access the first matrix by rows, but the second matrix is accessed by columns first, ie jumping by `B.size()`. If we can access matrix B row-wise we can utilize AVX instructions to perform packed floating point operations.

$$AB[i][j] = \sum_{k=0}^{A[0].size()-1} A[i][k] \times B[k][j]$$

We can observe that the order of the accessing order does not matter, as it does not change the way each cell in the result matrix is updated, therefore algebraically the result is the same; numerically, floating-point non-associativity may cause small rounding differences, but that's a trade-off for better performance.

Original

```

1 for (i = 0; i < size; i++)
2   for (j = 0; j < size; j++)
3     for (k = 0; k < size; k++)
4       AB[i][j] += A[i][k] * B[k][j];





```

Optimized

```

1 for (i = 0; i < size; i++)
2   for (k = 0; k < size; k++)
3     for (j = 0; j < size; j++)
4       AB[i][j] += A[i][k] * B[k][j];

```



## Reproduce result

compile & perf sampling options:

```

g++ mmomp.cpp -o mmomp -fopenmp -O3
srun perf stat -r 3 -e fp_arith_inst_retired.scalar_single,
    fp_arith_inst_retired.128b_packed_single, fp_arith_inst_retired.256b_packed_single,
    cycles,instructions ./mmomp
g++ rowbase.cpp -o rowbase -fopenmp -O3
srun perf stat -r 3 -e fp_arith_inst_retired.scalar_single,
    fp_arith_inst_retired.128b_packed_single, fp_arith_inst_retired.256b_packed_single, cycles,
    instructions, ./rowbase

```